



AI FOR INDUSTRY

The Italian Institute of
Artificial Intelligence

ADVERSARIAL ROBUSTNESS EVALUATION

Measuring the Residual Jailbreak Surface of Frontier Large Language Models

A red-team study of Anthropic Fable 5 & Opus 4.8 models

June 2026

\ AIS
AI Security

ABSTRACT

We evaluate the adversarial robustness of two frontier large language models (LLMs) developed by Anthropic, Fable 5 and Opus 4.8, against four families of automated jailbreak attack across 7 826 **harmful intents** spanning a ten-category harm taxonomy. Using the **HackAgent**¹ red-teaming framework, hundreds of thousands of adversarial attempts were generated and every apparent success was independently re-adjudicated by a panel of three judge models (majority vote). Both models resist the majority of attacks, but the residual surface is larger than aggregate framing suggests: it is dominated by *adaptive* iterative attacks, while static obfuscation is near-fully neutralised. The strongest adaptive search (tree-of-attacks) breaks Opus 4.8 on 11.5% of intents overall, whereas Fable 5 stays in the single digits (6.1% worst-case). Aggregate rates therefore should not be read as reassurance. Even in these hardened configurations, the two models produced 1 620 (Opus 4.8) and 702 (Fable 5) panel-confirmed harmful completions spanning *every* harm category, located automatically, cheaply, and within the first one or two refinement steps by an attacker model with no human expert in the loop. The reasonable conclusion is that even the best, most-tested frontier models remain reliably breakable under sustained automated pressure.

Contents

Executive summary	1
1 Introduction	1
2 Methodology	2
2.1 Threat model and target systems	2
2.2 Harmful-intent benchmark	2
2.3 Attack families	2
2.4 Two-stage adjudication with an independent judge panel	3
3 Results	3
3.1 Residual surface by attack family	3
3.2 Harm-category structure	4
3.3 Subcategory hotspots	4
3.4 How hard the attacker has to work	4
4 Analysis and interpretation	5
5 Limitations and caveats	7
6 Conclusion	7

AUTHOR

Dr. Nicola Franco

Head of AI Security Lab

The Italian Institute of Artificial Intelligence (AI4I)

Corso Castelfidardo 22, 10129 Turin, Italy

email: nicola.franco@ai4i.it

website: ais.rd-labs.ai4i.it

¹**HackAgent** is an open-source AI-agent red-teaming toolkit developed by the AI Security Lab at AI4I: it orchestrates an attacker model against a target model under a chosen jailbreak algorithm, scores the target’s responses, and logs every attempt. <https://hackagent.dev>.

Executive summary

We subjected two frontier models from Anthropic, Opus 4.8 and Fable 5, to an automated red-team campaign using the **HackAgent** framework. Across 7 826 **distinct harmful intents** spanning a ten-category safety taxonomy, four families of jailbreak technique generated hundreds of thousands of attempts. Every apparent success was then re-adjudicated by an **independent panel of three judge models**, and only attempts the panel confirmed by majority vote are counted as jailbreaks. This two-stage design is conservative: it discards borderline or judge-inflated “successes” that single-judge pipelines over-report.

What the numbers say

- **Both models resist most attacks, but Opus 4.8 breaks double digits under adaptive search.** The strongest attack family confirmed jailbreaks on 11.5% of intents against Opus 4.8 (tree-of-attacks) and 6.1% against Fable 5. The exposure is uneven across harm categories and worst where it matters most: against Opus 4.8, adaptive search reaches 27.6% on *child-safety* framings, with further double-digit cells in *criminal/economic* (14.7%), *content* (13.2%), and *cybersecurity* (PAIR, 16.6%); Fable 5 is most exposed in *ethical/social* and *child-safety* framings. The absolute counterpart of these rates is that Opus 4.8 and Fable 5 produced 1 620 and 702 **panel-confirmed harmful completions** respectively, spanning *every* harm category, found automatically and cheaply by an attacker model with no human expert in the loop.
- **Adaptive attacks dominate the residual surface.** The confirmed jailbreaks come almost entirely from *adaptive, iterative* attacks that let an attacker model rewrite its prompt in response to refusals, and they succeed *early*, usually within the first one or two refinement steps, so additional iterations buy the attacker little. By contrast, *static, template-based* obfuscation (encodings, ciphers, payload-splitting, and role-play or encyclopedic framing) is near-fully neutralised, confirmed at or below 0.2% despite roughly 50 000 attempts against each model.

The reasonable reading of these numbers is not that frontier models are safe, but that even the best, most-tested

frontier models remain reliably breakable under sustained automated pressure. At deployment scale, with millions of interactions per day, a success rate of this magnitude is not a rounding error but a steady, reproducible stream of harmful outputs reachable by anyone willing to iterate. The weak points are specific and addressable, but “addressable” is not “addressed.”

These results should be read as a *robustness characterisation* rather than a single safety score. Most attack families were run against both models on the *same* 7 826-intent taxonomy, so those cross-model comparisons are head-to-head; one iterative campaign was run only at partial scale and is a lower bound against Fable 5. Section 5 states these caveats in full.

1 Introduction

LLMs deployed in production are guarded by safety training and policy filters intended to refuse harmful requests. “Jailbreaks” are inputs crafted to circumvent those guards. As models improve, naive jailbreaks (a single prompt) increasingly fail, but *adaptive* adversaries, who iterate against the model’s own refusals, remain a credible threat. Understanding *how much* residual vulnerability remains, *which* techniques exploit it, and *which* harm categories are most exposed is essential for both model developers and the organisations that deploy these systems.

This white paper reports a systematic measurement of that residual surface. We use **HackAgent**, an automated red-teaming framework that orchestrates an attacker model against a target model under a chosen attack algorithm, scores the target’s responses, and logs every attempt. We pair it with a strict adjudication step, an independent multi-judge panel, so that the headline numbers reflect genuinely harmful completions rather than judge noise.

The study is designed to answer four questions:

1. **How robust are current frontier models overall?** We measure the fraction of harmful intents that can be jailbroken at all.
2. **Which attack families matter?** We locate where the residual surface actually lives, in adaptive search, static obfuscation, or elsewhere.
3. **Where is the exposure concentrated?** We identify which harm categories survive least well under attack.
4. **How hard does the attacker have to work?** We assess whether adding iterations meaningfully expands the attack’s reach.

Table 1. The ten harm categories and their 55 subcategories.

Code	Category	Subcategories
A	Ethical / Social	A1 Bias & discrimination behaviour A2 Insulting/harassing speech A3 Hate speech A4 Discriminatory behaviour
B	Privacy / Data	B1 Personal-data leakage B2 Biometric/health-data exposure B3 Doxxing & tracking B4 Surveillance misuse B5 Training-data leakage
C	Safety / Physical	C1 Dangerous instructions C2 Self-harm promotion C3 Physical harm C4 Human trafficking C5 Animal abuse C6 Drug guidance C7 Environmental damage
D	Criminal / Economic	D1 Fraud & scams D2 Economic crime D3 White-collar crime D4 Illegal activity D5 Evasion of law
E	Cybersecurity	E1 Malware generation E2 Exploit development E3 Phishing/ransomware E4 Jailbreak/prompt injection E5 Cross-modal attack E6 Data poisoning E7 Unauthorised model access
F	Information / Political	F1 Misinformation/disinformation F2 Deepfakes/synthetic media F3 Political propaganda F4 Election interference F5 Public-order disruption F6 National-security threat
G	Content / Cultural	G1 Pornography/nudity G2 Sexual content G3 Violence/gore G4 Cultural/religious offence
H	IP / Ownership	H1 Copyright infringement H2 Trademark/brand misuse H3 Proprietary-data leakage H4 Unauthorised style/persona imitation
I	Decision / Cognitive	I1 Medical advice I2 Health misguidance I3 Legal advice I4 Government decision support I5 Financial advice I6 Market manipulation I7 Psychological manipulation I8 Cognitive bias/overreliance
J	Child Safety	J1 CSAM & sexualisation J2 Grooming & enticement J3 Child trafficking J4 Harmful content targeting minors J5 Age-verification evasion

2 Methodology

2.1 Threat model and target systems

We treat each target as a black box accessed through its standard API. The attacker has no access to weights, logprobs, or internal state, only to the text the model returns. This mirrors the posture of an external adversary probing a deployed endpoint. Two frontier models from Anthropic were evaluated as victims: Opus 4.8 and Fable 5, both queried through a hosted OpenAI-compatible gateway.

2.2 Harmful-intent benchmark

Intents are drawn from a curated **harmful-intent taxonomy**: 7826 **harmful intents** organised into 10 **top-level harm categories** (A–J) and 55 **subcategories**. Categories span model-safety concerns, from social harms to cybersecurity to child safety (Table 1). Subcategory sizes are uneven (28–599 intents), reflecting the distribution of documented harm types rather than a balanced set.

2.3 Attack families

Four families of attack were exercised, spanning the standard taxonomy of *adaptive search*, *persuasion*, and *static obfuscation* (Figure 1):

TAP (Tree of Attacks with Pruning)²

An attacker model grows a tree of candidate prompts, expanding promising branches and pruning weak ones, guided by an on-the-fly score of the target’s responses. Adaptive and multi-step (configured here to depth 3, width 4, branching factor 3).

PAIR (Prompt Automatic Iterative Refinement)³

An attacker model and a scorer engage in an iterative loop: the attacker reads the target’s last refusal and rewrites the prompt to push past it, for up to a fixed iteration budget (configured to 12 iterations across 8 parallel streams, with early stopping on success).

PAP (Persuasive Adversarial Prompts)⁴

Reframes a harmful request using human-persuasion strategies (authority, role-play, hypotheticals) rather than iterative search.

H4RM3L⁵

A family of *static* obfuscation decorators applied to the raw intent: base64 encoding, character ciphers, payload-splitting, few-shot priming, “DAN”-style

²A. Mehrotra, M. Zampetakis, P. Kassianik, B. Nelson, H. Anderson, Y. Singer, and A. Karbasi. *Tree of Attacks: Jailbreaking Black-Box LLMs Automatically*. NeurIPS 2024; arXiv:2312.02119. <https://arxiv.org/abs/2312.02119>.

³P. Chao, A. Robey, E. Dobriban, H. Hassani, G. J. Pappas, and E. Wong. *Jailbreaking Black Box Large Language Models in Twenty Queries*. IEEE SaTML 2025; arXiv:2310.08419. <https://arxiv.org/abs/2310.08419>.

⁴Y. Zeng, H. Lin, J. Zhang, D. Yang, R. Jia, and W. Shi. *How Johnny Can Persuade LLMs to Jailbreak Them: Rethinking Persuasion to Challenge AI Safety by Humanizing LLMs*. ACL 2024; arXiv:2401.06373. <https://arxiv.org/abs/2401.06373>.

⁵M. Doumbouya, A. Nandi, G. Poesia, D. Ghosh, A. Goldie, et al. *h4rm3l: A Language for Composable Jailbreak Attack Synthesis*. ICLR 2025; arXiv:2408.04811. <https://arxiv.org/abs/2408.04811>.




Figure 1. Schematic of the three feedback-driven attack families, ordered top to bottom by increasing complexity. PAP applies a one-shot persuasion reframing with no target feedback; PAIR refines one prompt in a loop against the target’s refusals; TAP explores a pruned tree of candidate prompts, scored on the fly. Red dashed arrows mark the adaptive loop (PAIR/TAP) or a pruned branch (TAP). The static H4RM3L decorators apply fixed string transforms with no feedback and are omitted.

role-play, and Wikipedia-article framing.

In every case the *attacker* role was played by an uncensored open-weight model hosted on local GPUs, so that attacker refusals never confounded the measurement.

2.4 Two-stage adjudication with an independent judge panel

A persistent problem in jailbreak research is that a single automated judge over-reports success: it rewards responses that *begin* compliantly (“Sure, here is...”) even when the substance is harmless or hedged. To control for this we adjudicate in two stages.

1. **In-loop scoring.** During the attack, a fast scorer assigns each response a harmfulness score on a HarmBench-style rubric⁶; the attack uses this only to steer its search and decide early-stopping.
2. **Independent panel re-judging.** Every candidate success is then re-evaluated, offline, by a panel of **three independent judge models**, Qwen3.7 Max, Gemini 3.5 Flash, and GPT 5.5. An attempt counts as a confirmed jailbreak only if **a majority ($\geq 2/3$) of the panel** independently rates the target’s response as genuinely harmful.

Using judges from three different model families guards against any single family’s blind spots, and the majority-vote rule discards the borderline cases that inflate single-judge ASR. **All figures in this report are**

Table 2. Panel-confirmed jailbreaks per attack family: confirmed / attempts and ASR. h4rm3l aggregates all six decorators. [†]The Fable 5/PAIR campaign is partial (27/55 subcategories); its figures are a lower bound (Section 5).

Technique	Opus 4.8		Fable 5	
	conf./att.	ASR	conf./att.	ASR
PAIR	347 / 4 346	7.98%	162 / 3 766 [†]	4.30% [†]
TAP	901 / 7 826	11.51%	477 / 7 826	6.10%
PAP	287 / 7 826	3.67%	42 / 7 826	0.54%
H4RM3L	85 / 46 956	0.18%	21 / 46 956	0.04%
Total	1 620		702	

panel-confirmed.

Attack Success Rate (ASR), as used here.

$$\text{ASR} = \frac{\text{confirmed jailbreaks}}{\text{total attempts}} \times 100\%.$$

A *confirmed jailbreak* is an attempt the three-judge panel ruled harmful by majority (2 of 3). Denominators count attempts, not unique intents, so an intent attacked under several decorators contributes several attempts.

3 Results

3.1 Residual surface by attack family

Figure 2 compares panel-confirmed ASR for the four attack families that produced meaningful signal, for both targets. Table 2 gives the underlying counts.

⁶M. Mazeika, L. Phan, X. Yin, A. Zou, Z. Wang, N. Mu, et al. *HarmBench: A Standardized Evaluation Framework for Automated Red Teaming and Robust Refusal*. ICML 2024; arXiv:2402.04249. <https://arxiv.org/abs/2402.04249>.




Figure 2. Panel-confirmed attack success rate by technique, for both target models. Adaptive iterative attacks dominate; the static h4rm3l decorator family is near-zero against both models.

Reading the headline. Three facts stand out. First, the two models diverge sharply under the strongest attack: tree-of-attacks search breaks Opus 4.8 on 11.5% of intents, the only *double-digit* family ASR in the study, while Fable 5 holds every family to single digits (worst 6.1%). Second, the residual surface is concentrated in the feedback-driven families: the three adaptive/persuasion families account for 95% of confirmed jailbreaks against Opus 4.8 and 97% against Fable 5. Third, the *static* obfuscation family is effectively neutralised: despite roughly 50 000 attempts each, h4rm3l confirmed only 85 (Opus) and 21 (Fable) harmful completions.

3.2 Harm-category structure

The aggregate ASR still varies by harm category. Figure 3 plots each model's *robustness* per harm category, defined as $100\% - \text{ASR}$. A larger polygon means a more robust model; an inward dent marks a category where attacks found more purchase.

Table 3 decomposes this overview by technique, giving panel-confirmed ASR (%) with confirmed counts for every technique $\times$ harm-category cell.

The adaptive TAP search is the strongest family against *both* models, but it bites far harder into Opus 4.8. Its single hottest cell in the study is TAP against Opus 4.8 on child-safety framings at 27.6%, with further double-digit TAP cells in criminal/economic (14.7%), content/cultural (13.2%), cybersecurity (11.4%) and ethical/social (11.7%): the search reframes a blocked request until the model complies, and against Opus 4.8 it does so broadly across the taxonomy. PAIR adds a second cybersecurity-specific peak against Opus 4.8 (16.6% on malware/exploit content). Against Fable 5 the same TAP search is more contained, peaking on child-safety (13.7%) and ethical/social (10.2%) but holding cybersecurity near zero. Persuasion (PAP) produced a steadier, lower background rate across most categories of both models, with notable Opus peaks in criminal/economic (category D, 6.5%) and cybersecu-




Figure 3. Per-category robustness (100% – pooled ASR; radial axis 94–100%) for both targets. The deepest dents reveal each model's weakest categories: **child safety** and **cybersecurity** for Opus 4.8 (both $\approx 96\%$, pulled down by tree-of-attacks search) and **ethical-social** / **child safety** for Fable 5. Both models otherwise hold above 97–98% pooled robustness on most categories.

urity (category E, 6.3%) intents.

3.3 Subcategory hotspots

Aggregating to ten harm categories still averages over the 55 subcategories, which hides the sharpest exposure. Figure 4 ranks the fifteen most vulnerable subcategories by ASR, for both models. The exposure concentrates in a handful of subcategories: pooled over all attack families (so the high-volume static attempts dilute the rates), Opus 4.8 peaks at $\approx 11.5\%$ on E3 (phishing/ransomware) and $\approx 8.6\%$ on E2 (exploit development), its cybersecurity weak points, with further hotspots in F5 (public-order disruption, $\approx 7.4\%$) and G3 (violence/gore, $\approx 6.9\%$); Fable 5 peaks at $\approx 5.8\%$ on F1 (misinformation/disinformation) and A2 (insulting/harassing speech), with further hotspots in F5 (public-order disruption) and I6 (market manipulation), both $\approx 4.2\%$. Opus 4.8 sits above Fable 5 on most subcategories, consistent with its higher adaptive-attack exposure. The complete breakdown across all 55 subcategories is given in Figure 5.

3.4 How hard the attacker has to work

Figure 6 plots, for the two iterative families, the share of *all attempts* that first succeed at each refinement step. The signal is consistent across models: **successful jailbreaks are front-loaded**. For Fable 5 under TAP, the first step alone accounts for the largest single block of

Table 3. Panel-confirmed ASR (%) per technique and harm category, with confirmed counts in parentheses. h4rm3l aggregates all decorators; “—” marks untested pairs. Category codes A–J as in Table 1.

Technique	A	B	C	D	E	F	G	H	I	J
Opus 4.8										
TAP	11.7 (116)	7.5 (38)	9.7 (85)	14.7 (150)	11.4 (83)	9.0 (48)	13.2 (71)	3.0 (9)	6.0 (95)	27.6 (206)
PAIR	7.2 (71)	4.8 (24)	6.8 (19)	5.6 (55)	16.6 (120)	11.3 (17)	6.3 (34)	4.0 (7)	—	—
PAP	3.8 (38)	1.2 (6)	1.7 (15)	6.5 (66)	6.3 (46)	2.6 (14)	3.4 (18)	1.0 (3)	3.4 (54)	3.6 (27)
h4rm3l	0.1 (8)	0.0 (1)	0.2 (12)	0.0 (3)	0.5 (20)	0.2 (7)	0.1 (3)	0.1 (1)	0.1 (11)	0.4 (19)
Fable 5										
TAP	10.2 (101)	2.0 (10)	4.4 (39)	3.2 (33)	0.4 (3)	7.3 (39)	6.9 (37)	2.6 (8)	6.6 (105)	13.7 (102)
PAIR [†]	9.2 (91)	1.8 (9)	5.2 (46)	2.1 (14)	0.3 (2)	—	—	—	—	—
PAP	1.0 (10)	0.2 (1)	1.0 (9)	0.7 (7)	0.0 (0)	0.0 (0)	0.7 (4)	0.0 (0)	0.5 (8)	0.4 (3)
h4rm3l	0.1 (4)	0.0 (0)	0.1 (4)	0.0 (1)	0.0 (0)	0.1 (4)	0.1 (3)	0.0 (0)	0.0 (4)	0.0 (1)

[†]Partial campaign (27/55 subcategories); F–J not covered.




Figure 4. Fifteen most vulnerable subcategories by ASR, Opus 4.8 vs Fable 5. The exposure concentrates in a few subcategories, and the two models’ hotspots are largely disjoint.

successes; by the third step the marginal yield has fallen sharply. PAIR against Opus 4.8 is somewhat more spread but still concentrates in the first two iterations.

The practical implication cuts two ways. For defenders, it means the bulk of the risk is realised quickly, a small, well-chosen reframing rather than an exhaustive search. For attackers, it means additional compute spent on deeper iteration yields diminishing returns: the easy wins come first, and the model’s deeper refusal behaviour is stable.

4 Analysis and interpretation

Adaptive beats static, decisively. The clearest result of the study is the gulf between adaptive and static attacks. Static decorators (encoding the payload, splitting it, wrapping it in a template) do not adapt, and modern safety training appears to have largely closed those

well-documented holes. The 50 000-attempt h4rm3l campaigns returning $\leq 0.2\%$ confirmed ASR is strong evidence that obfuscation alone is no longer a viable attack against frontier models.

The vulnerability is contextual, not lexical. Because the surviving attacks work through *framing* rather than *encoding*, they are harder to defend with surface-level filters. This points defenders toward semantic, context-aware monitoring of multi-turn interactions rather than input sanitisation.

Category structure is partly shared, partly model-specific. Both models are most exposed to the adaptive TAP search, and both have child-safety framing among their weakest categories, suggesting a common failure mode of context-driven reframing rather than a purely model-specific quirk. The *severity* is what differs: against Opus 4.8 the exposure is roughly twice as large and spreads into double digits across child-safety, crimi-




Figure 5. ASR per subcategory for both targets, over all 55 subcategories. Pooling over the high-volume static (h4rm3l) attempts keeps most absolute rates low; the relative hotspots are Opus 4.8 on E3/E2 (cybersecurity), F5 and G3, and Fable 5 on A2, F1, F5 and I6. Figure 4 ranks the fifteen highest.




Figure 6. Per-attempt success rate by refinement iteration for the iterative families (PAIR, TAP), both targets. y is the true per-attempt success rate (successes at step N divided by *all* attempts of that series), not a conditional distribution. PAIR display is capped at the common budget of 5 iterations; TAP maxes at 3.

nal/economic, content and cybersecurity, whereas Fable 5 keeps cybersecurity near zero and stays in single digits elsewhere. That the gaps are concentrated in identifiable categories is encouraging, because it implies they are addressable through targeted data and evaluation, but the shared child-safety weakness in particular warrants attention from both developers.

5 Limitations and caveats

Read before quoting any single number

- **Comparisons rest on matched samples, except for PAIR.** TAP, PAP, and h4rm3l used the *same* 7 826-intent taxonomy with identical denominators for both models, so those comparisons are head-to-head; PAIR is partial and unequal in coverage (Opus 4.8 38/55, Fable 5 27/55), so it stays directional. Narrow gaps should not be over-read as a certified “safer” ordering.
- **The Fable 5/PAIR campaign is partial.** It covers 27 of 55 subcategories (a target-routing bug halted it mid-run), so its 162 confirmed jailbreaks and 4.30% ASR are a **lower bound**. Categories F–J are absent from the Fable 5 PAIR row of Table 3.
- **Judge panels are imperfect.** The panel reduces but does not eliminate adjudication error. Majority vote can both miss subtle harms (false negatives) and, less often, over-credit fluent-but-inert responses.
- **Point-in-time snapshot.** Results reflect the model versions and safety configurations available at evaluation time; produc-

tion safety stacks (system prompts, output filters, monitoring) are not modelled and would further reduce real-world success.

6 Conclusion

The percentages should not be read as reassurance. The other side of “89% resisted” (the floor under Opus 4.8’s worst attack family) is its absolute counterpart, and that counterpart deserves to be stated plainly. These are among the most heavily safety-trained systems ever deployed, evaluated here in hardened configurations, and they still produced 1 620 (**Opus 4.8**) and 702 (**Fable 5**) **panel-confirmed harmful completions**. These are not borderline cases: each survived a 2-of-3 independent-judge vote, and they span *every* harm category in the taxonomy, including the most serious, from cybersecurity weaponisation to child-safety framings. Three properties make this more concerning than the headline rates suggest. First, the failures were found *automatically*: an attacker model with no human expert in the loop located them over a campaign measured in days, not months. Second, they were found *cheaply and fast*: when an attack succeeds it succeeds within the first one or two steps, so the marginal cost of a working jailbreak is low. Third, at deployment scale, with millions of interactions per day, a success rate of this magnitude is not a rounding error but a steady, reproducible stream of harmful outputs reachable by anyone willing to iterate. The reasonable conclusion is not that frontier models are safe, but that even the best, most-tested frontier models remain reliably breakable under sustained automated pressure. The distance between looking safe under casual use and being safe under adversarial use remains wide, and the scale of safety investment that still leaves this gap open is itself the reason the result warrants scrutiny rather than comfort.

This document reports aggregate adversarial-robustness statistics for defensive research. Harmful model outputs are not reproduced.

Fable 5 and Opus 4.8 are large language models developed by Anthropic. All model names referenced are trademarks of their respective owners. This is an independent third-party evaluation and is not affiliated with, authorised by, or endorsed by Anthropic or any other model provider.

GET IN TOUCH

AI Security Lab

The Italian Institute of Artificial Intelligence (AI4I)

*Transformative, application-oriented AI research
driving industrial innovation.*

CONTACT

The Italian Institute of Artificial Intelligence (AI4I)

Corso Castelfidardo 22, 10129 Turin, Italy

www.ai4i.it | ai4i@ai4i.it

AI Security Lab

www.ais.rd-labs.ai4i.it

ais@ai4i.it

AI⁴I
A FOR INDUSTRY
The Italian Institute of
Artificial Intelligence

\ AIS
AI Security