

MUST KNOW

RAG

Interview Questions

12

SECTIONS

69

QUESTIONS

∞

VALUE



Chandra Sekhar

AI Engineer & Trainer

Bundle link in **Caption - Enroll Now**

12 Topics · 69 Questions



RAG Fundamentals & Core Concepts

6 questions



RAG Architecture & Workflow

5 questions



Document Loading & Data Ingestion

7 questions



Text Splitting & Chunking Strategies

9 questions



Embeddings & Vector Stores

5 questions



Vector Search & Retrieval Techniques

8 questions



Vector Store Operations & Debugging

3 questions



Prompting & Context Injection

5 questions



Evaluation, Monitoring & Observability

6 questions



Performance, Cost & Scalability

5 questions



Security, Privacy & Enterprise Readiness

5 questions



Advanced & Agentic RAG (Bonus)

5 questions

Every question here has been asked in real AI Engineer interviews.



SECTION 1

RAG Fundamentals & Core Concepts

- 1 What is Retrieval Augmented Generation (RAG), and why is it critical for LLMs?
- 2 What are the limitations of a pure LLM system without RAG?
- 3 How is RAG different from fine-tuning an LLM?
- 4 When should you prefer RAG over fine-tuning?
- 5 What types of problems are best solved using RAG?
- 6 What kinds of data sources can be used in a RAG system?

Bundle link in **Caption** - Enroll Now



SECTION 2

RAG Architecture & Workflow

- 7 What are the two main components of a RAG application architecture?
- 8 How does the indexing process work, and why is it essential?
- 9 What is the step-by-step workflow of a complete RAG pipeline?
- 10 What is the specific role of the Retrieval component?
- 11 What is the goal of integrating retrieval with generation?

Bundle link in **Caption** - Enroll Now



SECTION 3

Document Loading & Data Ingestion

- 12 What is the primary purpose of DocumentLoaders?
- 13 How does WebBaseLoader function within the LangChain ecosystem?
- 14 What are the roles of WebBaseLoader and RecursiveCharacterTextSplitter?
- 15 How can you customize the HTML parsing process in WebBaseLoader?
- 16 What is the specific role of BeautifulSoup in document ingestion?
- 17 What is a SoupStrainer and how does it improve efficiency?
- 18 Why is customized HTML parsing critical for a RAG system?

Bundle link in **Caption** - Enroll Now



SECTION 4

Text Splitting & Chunking Strategies

- 19 Why is it necessary to split documents into smaller chunks?
- 20 Why must long documents be split before embedding?
- 21 What is chunk overlap and why is it used?
- 22 How does chunk size impact retrieval accuracy?
- 23 What trade-offs exist between larger chunks and smaller chunks?
- 24 How does poor chunking lead to hallucinations?
- 25 Why is RecursiveCharacterTextSplitter recommended for generic text?
- 26 How does the `'add_start_index=True'` parameter

Bundle link in **Caption** - Enroll Now



SECTION 5

Embeddings & Vector Stores

- 28** What is the purpose of embedding text and storing it in a vector store?
- 29** How does the choice of embedding model affect RAG performance?
- 30** What happens if different embedding models are used for indexing and querying?
- 31** What are the key components needed to build and query a vector store?
- 32** What is Cosine Similarity and why is it used in vector searches?

Bundle link in **Caption** - Enroll Now

SECTION 6



Vector Search & Retrieval Techniques

- 33** What is the specific role of a Retriever in LangChain?
- 34** How does the VectorStoreRetriever function?
- 35** What does the `k` parameter control in vector retrieval?
- 36** What is the difference between similarity search and MMR (Max Marginal Relevance)?
- 37** What is hybrid search and why is it useful in RAG systems?
- 38** How do metadata filters improve retrieval quality?
- 39** What is reranking and why is it important in advanced RAG pipelines?

Bundle link in **Caption** - Enroll Now



SECTION 7

Vector Store Operations & Debugging

- 41** What are common challenges when working with the Chroma vector store?
- 42** How do you inspect the contents of a vector store?
- 43** What is the step-by-step workflow to embed and store documents?

Bundle link in **Caption** - Enroll Now



SECTION 8

Prompting & Context Injection

- 44 How is retrieved context injected into the prompt?
- 45 What is context stuffing and why is it dangerous?
- 46 How do you design prompts to prevent hallucinations in RAG?
- 47 What happens if retrieved documents contradict each other?
- 48 How do you limit the LLM to answer only from retrieved context?

Bundle link in **Caption** - Enroll Now

SECTION 9



Evaluation, Monitoring & Observability

- 49** How do you evaluate the quality of a RAG system?
- 50** What metrics are used to measure retrieval quality?
- 51** How do you detect hallucinations in a RAG pipeline?
- 52** Why is configuring LangSmith crucial for complex applications?
- 53** How does LangSmith help in debugging RAG applications?
- 54** What are common failure modes in RAG systems?

Bundle link in **Caption** - Enroll Now



SECTION 10

Performance, Cost & Scalability

- 55 How do you reduce latency in a RAG system?
- 56 How do you optimize embedding costs?
- 57 What caching strategies are used in RAG pipelines?
- 58 How do you handle large-scale document updates?
- 59 How do you design a RAG system for production use?

Bundle link in **Caption** - Enroll Now

SECTION 11

Security, Privacy & Enterprise Readiness

- 60** How does RAG help with private or sensitive data?
- 61** What are the security risks in a RAG-based system?
- 62** How do you prevent data leakage in RAG responses?
- 63** How do access controls work in enterprise RAG systems?
- 64** How do you ensure compliance when using RAG?

Bundle link in **Caption** - Enroll Now



SECTION 12

Advanced & Agentic RAG (Bonus)

- 65 What is Agentic RAG?
- 66 How does RAG work with tools and agents?
- 67 What is self-reflection or self-correction in RAG systems?
- 68 How does multi-step retrieval improve answer quality?
- 69 What is RAG with feedback loops?

Bundle link in **Caption** - Enroll Now



Want Access to the Entire AI Interview Bundle?

13

COURSES

700+

QUESTIONS

Interview

READY

 INTRODUCING

The AI Interview Master Bundle

13 Interview Courses 650+ Questions



13 Structured Interview Courses

Topic-wise, end-to-end prep



650+ Carefully Curated Questions


No fluff — interview-ready



Beginner → Advanced

One structured path

Check Next slide for more details



Overview of AI Interview Mastery Bundle

13 modules • 675+ questions

1

Python For GenAI

55 questions

2

Statistics & Math for ML

30 questions

3

Machine Learning Fundam...

70 questions

4

Deep Learning Basics

25 questions

5

LLM Core Concepts

50 questions

6

RAG Interview Mastery

69 questions

7

LangChain Core Concepts

134 questions

8

LangGraph Core Concepts

49 questions

9

Agentic AI

50 questions

10

CrewAI

53 questions

11

AutoGen

28 questions

12

FastAPI (Backend for AI)

32 questions

13

Docker for AI Engineers

30 questions

• Access bundle link from **caption**

GET HIRED

The only bundle you'll ever need.



Stop scattered prep. Start structured learning.

13 modules · 645 questions · built for AI engineers



Link of Bundle in caption



Chandra Sekhar

AI Engineer and Trainer

Follow for daily AI interview
questions, tips & career advice.

● [Follow on LinkedIn →](#)