

TUTORIAL

Ollama

The Complete Setup Guide

```
$ ollama pull qwen3:30b
pulling manifest...
pulling layers... [████████████████████████████████████████] 100%
success
```

```
$ ollama run qwen3:30b
>>> Hello! How can I help you today?
```

↑ two commands,
that's it

THE BASICS

What is Ollama?

Think of it as **Docker for AI models**.

Pull any model with one command.
It handles quantization, memory, and GPU acceleration automatically.

Mac. Windows. Linux.


500+ models. Zero cloud accounts.



100%
Private



Zero
API Costs




Full
Control

STEP 01


Install Ollama

Download from ollama.com or run one command.



Option A: Download

OR



Option B: One command

works on Mac, Windows, and Linux

STEP 02

Pull Your First Model

One command downloads
any AI model. Ollama
handles the rest.

```
$ ollama pull qwen3:30b
```

```
pulling manifest
```

```
pulling sha256:a4b6...
```

```
[████████████████████████████████████████] 100% 18 GB/18 GB
```

```
success
```

← swap this
for any model

llama3.3

deepseek-r1

qwen3

gemma4

500+ more at ollama.com/library

STEP 03

Start Chatting

That's it. You're talking to an AI running entirely on your machine.

```
$ ollama run qwen3:30b
```

```
>>> Explain Docker in 3 sentences
```

Docker packages your app and its dependencies into a container. That container runs the same way on any machine. Think of it as a lightweight, portable box for your code.

```
>>> _
```

running
100%
locally

no API key, no internet needed

TOP PICKS

Best Models to Start With



Qwen 3 30B

Best all-rounder.
Chat, code, reasoning.

18 GB



DeepSeek R1 14B

Thinks before answering.
Great for logic.

9 GB



Llama 3.3 8B

Fast and light.
Runs on anything.

5 GB



Gemma 4 12B

Strong for image
understanding.


8 GB

size = how much RAM/VRAM you need

LEVEL UP

Connect It to Everything

Ollama runs a local API. Plug it into any tool.



WHY

Why Go Local?



100% Private

Your data never leaves your machine.




Zero API Costs

Run it as much as you want, forever.



Works Offline

No internet required after download.



Full Control

Pick your model. Tweak everything. No limits.

~~ChatGPT API:
~\$20/month~~

Ollama: \$0 ✓